


Lecture 8: Logistic Regression

Andreas Wichert

Department of Computer Science and Engineering
Técnico Lisboa

- In statistics the artificial neuron with sigmoid activation function is called logistic regression
- The sigmoid function called the logistic function.



- Suppose we want to predict whether someone is female C_1 or male C_2 using height in centimeters
- If linear regression is used $net = y(\mathbf{w}, x) = \mathbf{w}^T \cdot \mathbf{x}$
- The predicted values will become greater than one and less than zero.
- Such values are inadmissible!

- However with

$$p(C_1|\mathbf{x}) = \frac{e^{(net)}}{1 + e^{(net)}} = \frac{1}{1 + e^{(-net)}} \in [0, 1]$$

$$p(C_2|\mathbf{x}) = 1 - p(C_1|\mathbf{x}) = \frac{1}{1 + e^{(net)}} \in [0, 1]$$

with odds being

$$odds = \frac{p(C_1|\mathbf{x})}{p(C_2|\mathbf{x})} = \frac{p(C_1|\mathbf{x})}{1 - p(C_1|\mathbf{x})} = e^{(net)}.$$


we can overcome this problem.

The *logit* function that is the inverse of logistic (sigmoid σ) function with

$$\log(odds) = \textit{logit}(\mathbf{w}^T \cdot \mathbf{x}) = \log \left(\frac{p(C_1|\mathbf{x})}{1 - p(C_1|\mathbf{x})} \right) = \mathbf{w}^T \cdot \mathbf{x} = \textit{net}$$



corresponds to the linear regression. It is common in machine learning community to call a sigmoid unit with cross entropy loss function a logistic regression.

Regression of Sex on Height



- None of the observations (data points) fall on the regression line
- They are all zero or one

Linear vs. Logistic regression



It follows

$$o = \sigma(net) = \sigma \left(\sum_{j=0}^N w_j \cdot x_j \right) = \sigma (\mathbf{w}^T \cdot \mathbf{x})$$

$$\sigma'(net) = \sigma(net) \cdot (1 - \sigma(net))$$

and

$$\sigma(-net) = 1 - \sigma(net)$$

We can think of the sigmoid output unit as using the linear input net and using the sigmoid activation function to convert net into a probability. It can be interpreted as logistic regression with two classes C_1 and C_2 with

$$p(C_1|\mathbf{x}) = \sigma \left(\sum_{j=0}^N w_j \cdot x_j \right) = \sigma (\mathbf{w}^T \cdot \mathbf{x})$$

and

$$p(C_2|\mathbf{x}) = 1 - p(C_1|\mathbf{x})$$

Training

How to train

$$Data = \{(\mathbf{x}_1, t_1), (\mathbf{x}_2, t_2), \dots, (\mathbf{x}_N, t_N)\}, \quad t_k \in \{0, 1\}$$

The likelihood function can now be written

$$p(\mathbf{t}|\mathbf{w}) = \prod_{k=1}^N p(C_1|\mathbf{x}_k)^{t_k} \cdot (1 - p(C_1|\mathbf{x}_k))^{1-t_k}$$

with

$$o_k = p(C_1|\mathbf{x}_k) = \sigma(\mathbf{w}^T \cdot \mathbf{x}_k) = \sigma(net_k)$$

$$p(\mathbf{t}|\mathbf{w}) = \prod_{k=1}^N o_k^{t_k} \cdot (1 - o_k)^{1-t_k}$$

Error function

- Error function is defined by negative logarithm of the likelihood

$$E(\mathbf{w}) = -\log(p(\mathbf{t}|\mathbf{w})) = -\sum_{k=1}^N (t_k \log o_k + (1 - t_k) \log(1 - o_k))$$

which gives the cross entropy error with

$$\neg t_k = (1 - t_k)$$

$$H(t, p) = -\sum_{k=1}^N (t_k \cdot \log(p(C_1|\mathbf{x}_k)) + \neg t_k \cdot \log(p(C_2|\mathbf{x}_k)))$$

$$\neg t_k = (1 - t_k)$$

$$H(t, p) = - \sum_{k=1}^N (t_k \cdot \log(p(C_1|\mathbf{x}_k)) + \neg t_k \cdot \log(p(C_2|\mathbf{x}_k)))$$

The gradient that minimizes the error (loss) function is given by

$$\frac{\partial E}{\partial w_j} = \frac{\partial}{\partial w_j} \left(- \sum_{k=1}^N (t_k \log o_k + (1 - t_k) \log(1 - o_k)) \right)$$

$$\frac{\partial E}{\partial w_j} = \frac{\partial}{\partial w_j} \left(- \sum_{k=1}^N (t_k \log o_k + (1 - t_k) \log(1 - o_k)) \right)$$

with

$$\frac{\partial E}{\partial w_j} = - \sum_{k=1}^N \left(\frac{t_k}{o_k} + - \frac{(1 - t_k)}{(1 - o_k)} \right) \cdot \frac{\partial}{\partial w_j}(o_k)$$

$$\frac{\partial E}{\partial w_j} = \sum_{k=1}^N \left(\frac{(1 - t_k)}{(1 - o_k)} - \frac{t_k}{o_k} \right) \cdot \frac{\partial}{\partial w_j}(o_k)$$

$$\frac{\partial E}{\partial w_j} = \sum_{k=1}^N \left(\frac{(1 - t_k)}{(1 - o_k)} - \frac{t_k}{o_k} \right) \cdot (o_k \cdot (1 - o_k)) \cdot \frac{\partial}{\partial w_j} \left(\sum_{j=0}^D w_j \cdot x_{k,j} \right)$$

$$\frac{\partial E}{\partial w_j} = \sum_{k=1}^N ((1 - t_k) \cdot (o_k)) - (t_k) \cdot (1 - o_k)) \cdot \frac{\partial}{\partial w_j} \left(\sum_{j=0}^D w_j \cdot x_{k,j} \right)$$


$$\frac{\partial E}{\partial w_j} = \sum_{k=1}^N (o_k - t_k) \cdot \frac{\partial}{\partial w_j} \left(\sum_{j=0}^D w_j \cdot x_{k,j} \right)$$

$$\frac{\partial E}{\partial w_j} = \sum_{k=1}^N (o_k - t_k) \cdot x_{k,j}$$

$$\frac{\partial E}{\partial w_j} = - \sum_{k=1}^N (t_k - o_k) \cdot x_{k,j}.$$

The update rule for gradient decent is given by

$$\Delta w_j = \eta \cdot \sum_{k=1}^N (t_k - o_k) \cdot x_{k,j}.$$



2 artificial neurons correspond to two perceptrons with the same input x and two target values represented by the vector $y = (y_1, y_2)$, the actual output computed by the network is $o = (o_1, o_2)$.

For K artificial neurons with a continuous activation function ϕ an index t is used to identify the corresponding artificial linear neuron with

$$t \in \{1, 2, \dots, K\}$$

to identify the weight vector \mathbf{w}_t and the t the output o_t from the neuron,


$$o_t = \phi \left(\sum_{j=1}^D w_{t,j} \cdot x_j \right)$$

Assuming that the training set consists of N observations

$$X = (\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_k, \dots, \mathbf{x}_N)^T$$

and respective target values represented as vectors of dimension K (since t is used as an index, we will use y_{kt} to indicate the specific target)

$$Y = (\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_k, \dots, \mathbf{y}_N)^T.$$

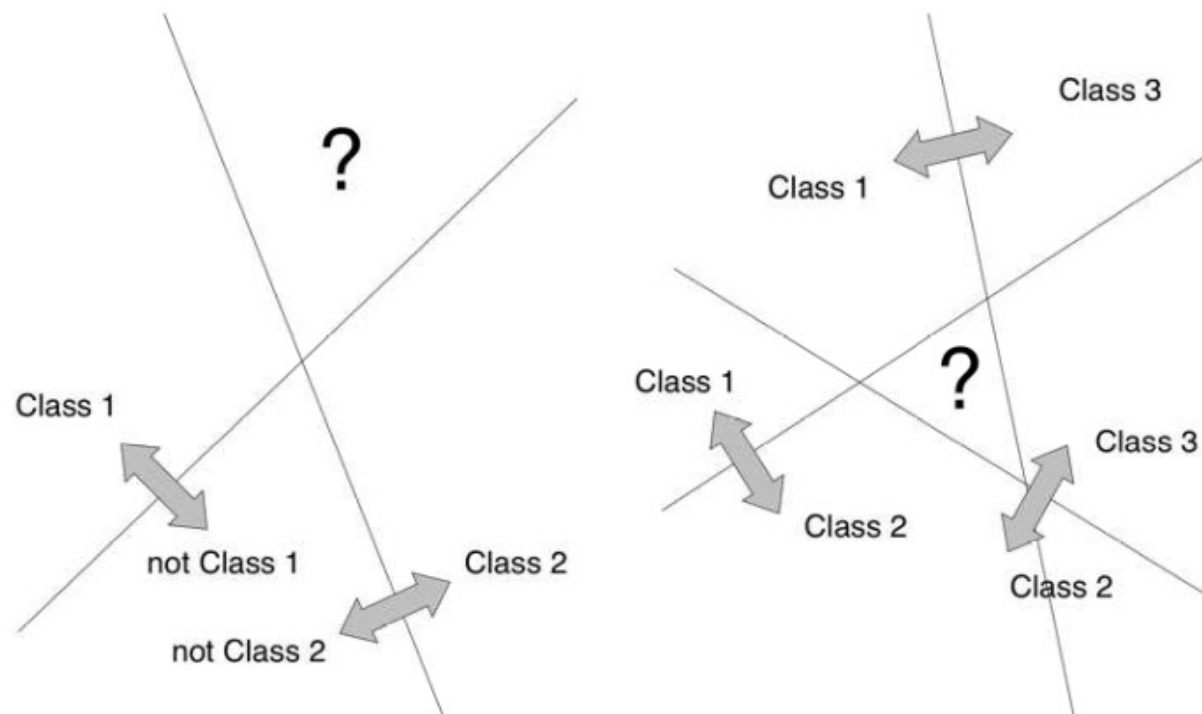


During training, each neuron is trained individually with its target value y_{kt}

$$y_{kt} \in \{0, 1\}, \quad \sum_{t=1}^K y_{kt} = 1$$

After training, the prediction for an input pattern \mathbf{x} is done using

$$\arg \max_t \left(\phi \left(\mathbf{w}_t^T \cdot \mathbf{x} \right) \right). \quad (5.51)$$



For the perceptron we cannot apply the $\arg \max_t$ operation since the output of each neuron is either 1 or 0 and building a $K - 1$ classifiers each of which solves a two class problem creates ambiguous regions

For K artificial neurons with the sigmoid activation the results cannot be interpreted as probabilities any more

$$o_t = \sigma \left(\sum_{j=1}^D w_{t,j} \cdot x_j \right) = \sigma(net_t)$$

since

$$\sum_{t=1}^K o_t \neq 1.$$

To be able to interpret them again as such, we can simply normalize the values with

$$o_s = \frac{\sigma(net_s)}{\sum_{t=1}^K \sigma(net_t)}.$$

Since each neuron can be viewed as a single classifier as well with odds

$$odds_s = \frac{p(C_{s,1}|\mathbf{x})}{p(C_{s,2}|\mathbf{x})} = \frac{p(C_{s,1}|\mathbf{x})}{1 - p(C_{1,s}|\mathbf{x})} = e^{(net_s)}.$$

$$\log(odds) = \text{logit}(\sigma(\mathbf{w}^T \cdot \mathbf{x})) = \log\left(\frac{p(C_1|\mathbf{x})}{1 - p(C_1|\mathbf{x})}\right) = \mathbf{w}^T \cdot \mathbf{x}$$

$$o_s = \frac{odd_s}{\sum_{t=1}^K odd_t} = \frac{\exp(net_s)}{\sum_{t=1}^K \exp(net_t)}$$

$$p(C_s|\mathbf{x}) = \sigma(\mathbf{net})_s = \frac{\exp(net_s)}{\sum_{t=1}^K \exp(net_t)}$$

$$p(C_\kappa|\mathbf{x}) = \sigma(\mathbf{net})_\kappa = \frac{p(\mathbf{x}|C_\kappa) \cdot p(C_\kappa)}{\sum_{j=1}^K p(\mathbf{x}|C_j) \cdot p(C_j)} = \frac{\exp(net_k)}{\sum_{j=1}^K \exp(net_j)}$$

- The **softmax function** is used in various multi class classification methods, such as multinomial logistic regression (also known as softmax regression) with the prediction

$$\arg \max_{\kappa} (\sigma(\mathbf{net})_\kappa) = \arg \max_{\kappa} (\mathbf{w}_\kappa^T \cdot \mathbf{x})$$

Cross Entropy Loss Function for softmax

- Taking the error of the gradient function with respect to the vector \mathbf{w}_j we obtain

$$\nabla_{\mathbf{w}_s} E(\mathbf{w}_1, \dots, \mathbf{w}_K) = \sum_{k=1}^N (o_{ks} - y_{ks}) \cdot \mathbf{x}_k = - \sum_{k=1}^N (y_{ks} - o_{ks}) \cdot \mathbf{x}_k$$

$$\sigma(net_{ks}) = \frac{\exp(net_{ks})}{\sum_{t=1}^K \exp(net_{kt})}$$

which is called the normalized exponential or softmax function. For simplification of the derivative function we drop the index k that indicates the k -th training pattern and we get

$$\frac{\partial \sigma(net_s)}{\partial net_t} = \sigma(net_s) \cdot (I_{st} - \sigma(net_t)) = o_s \cdot (I_{st} - o_t)$$

means if $t = s$

$$\frac{\partial \sigma(net_s)}{\partial net_t} = \sigma(net_s) \cdot (1 - \sigma(net_s)) = o_s \cdot (1 - o_s)$$

if $t \neq s$

$$\frac{\partial \sigma(net_s)}{\partial net_t} = -\sigma(net_t) \cdot \sigma(net_s) = -o_t \cdot o_s$$

Kronecker Function

with the Kronecker Function

$$I_{st} = \delta_{st} = \begin{cases} 1, & \text{if } s = t \\ 0, & \text{otherwise.} \end{cases}$$

- Taking the error of the gradient function with respect to net_j over all training patterns we get

$$E(\mathbf{w}) = - \sum_{k=1}^N \sum_{t=1}^K y_{kt} \cdot \log o_{kt}$$

$$\frac{\partial E}{\partial o_{kt}} = - \sum_{k=1}^N \sum_{t=1}^K \frac{y_{kt}}{o_{kt}}, \quad \frac{\partial o_{kt}}{\partial net_s} = o_{kt} \cdot (I_{ts} - o_{ks}),$$

$$\frac{\partial E}{\partial net_s} = - \sum_{k=1}^N \sum_{t=1}^K \frac{y_{kt}}{o_{kt}} \cdot o_{kt} \cdot (I_{ts} - o_{ks})$$

$$\frac{\partial E}{\partial net_s} = - \sum_{k=1}^N \sum_{t=1}^K (y_{kt} \cdot (I_{ts} - o_{ks})),$$

since

$$\sum_{t=1}^K y_{kt} = 1,$$

$$\frac{\partial E}{\partial net_s} = - \sum_{k=1}^N (y_{ks} - o_{ks})$$

Since

$$net_{ks} = \sum_{j=1}^D w_{js} x_{kj}$$

we get for all N training patterns

$$\frac{\partial E}{\partial w_{js}} = - \sum_{k=1}^N (y_{ks} - o_{ks}) \cdot x_{kj}$$

$$\frac{\partial E}{\partial w_{js}} = - \sum_{k=1}^N (y_{ks} - o_{ks}) \cdot x_{kj}$$

with the learning rule

$$\Delta w_{js} = \eta \cdot \sum_{k=1}^N (y_{ks} - o_{ks}) \cdot x_{kj}$$

$$w_{js}^{new} = w_{js}^{old} + \Delta w_{js}.$$

This algorithm is called the logistic regression.

Logistic Regression Algorithm

Given a training set (sample)

$$Data = \{(\mathbf{x}_1, \mathbf{y}_1), (\mathbf{x}_2, \mathbf{y}_2), \dots, (\mathbf{x}_k, \mathbf{y}_k), \dots, (\mathbf{x}_N, \mathbf{y}_N)\}$$

with \mathbf{y}_k represented as vectors of dimension K . During the training each neuron is trained individually with its target value y_{kt}

$$y_{kt} \in \{0, 1\}, \quad \sum_{t=1}^K y_{kt} = 1$$

the goal of the algorithm is to correctly classify the test set (population) into K classes $C_1 = 100\dots$, $C_2 = 010\dots$, $C_3 = 001\dots$, \dots

- (1) iterations=0;
- (2) $\eta \in (0, 1]$;
- (3) FOR t=1 TO K
 - {
 - Initialise all the weights $w_{0t}, w_{1t}, \dots, w_{Dt}$ to some random values;
 - }
- (4) Choose a pattern \mathbf{x}_k out of the training set;
- (5) FOR t=1 TO K
 - {
 - Compute $net_{kt} = \sum_{i=1}^D w_{jt} \cdot x_{kj} + w_{0t} = \langle \mathbf{x}_k | \mathbf{w}_t \rangle + w_{0t} \cdot x_0$;
 - Compute $odds_{kt} = \exp(net_{kt})$;
 - }
- (6) FOR t=1 TO K
 - {
 - Compute $o_{kt} = \frac{odds_{kt}}{\sum_t odds_{kt}}$;
 - Compute $\Delta w_{jt} = \eta \cdot (y_{kt} - o_{kt}) \cdot x_{k,j}$;
 - Update the weights $w_{jt} = w_{jt} + \Delta w_{jt}$;
 - }
- (7) iterations++;
- (8) If no change in weights for all training set or maximum number of iteration THEN STOP ELSE GOTO 4;

- An epoch corresponds to adapting all trainings vectors (patterns) \mathbf{x}_k with $k = 1, 2, \dots, N$, with $\text{epoch} = \text{iterations}/N$.
- Depending on the training set, the initialization and the size of the learning rate, the epochs have to be repeated several times until a solution is reached.

Sigmoid Unit versus Logistic Regression

Sigmoid Unit is with target, should be positive (between zero and one):

$$\Delta w_j = \eta \cdot \alpha \cdot \sum_{k=1}^N (t_k - o_k) \underbrace{\sigma(\text{net}_{k,j}) \cdot (1 - \sigma(\text{net}_{k,j}))}_{\text{Sigmoid Unit}} \cdot x_{k,j}$$

Logistic Regression is with target $t_k \in \{0, 1\}$

$$\Delta w_j = \eta \cdot \sum_{k=1}^N (t_k - o_k) \cdot x_{k,j}.$$

If we assume $\alpha = 1$ then the difference between sigmoid unit and the logistic regression that was derived by maximising the negative logarithm of the likelihood is

$$\sigma(\text{net}_{k,j}) \cdot (1 - \sigma(\text{net}_{k,j})) \geq 0$$

the step size in the direction of gradient. Does it mean that Sigmoid Unit converge faster?

Linear Unit versus Logistic Regression

Target can be any value and can be solved by closed-form solution, by pseudo inverse

$$o_k = \sum_{j=0}^D w_j \cdot x_{k,j}$$

Target $t_k \in \{0, 1\}$ cannot be solved by closed-form solution

$$o_k = \frac{1}{1 + e^{(-\alpha \cdot (\sum_{j=0}^N w_j \cdot x_{k,j}))}}$$

$$\Delta w_j = \eta \cdot \sum_{k=1}^N (t_k - o_k) \cdot x_{k,j}.$$

Logistic Regression as well as the sigmoid unit gives a better decision boundary.

For Sigmoid (Logistic) distant points from the decision boundary have the same impact







Figure 1.2: (a) Linear activation function. (b) The function $\sigma(\text{net})$ with $\alpha = 5$


Distance




Better Decision Boundary of Logistic Regression LR (sigmoid) to Linear Unit




Literature



- Simon O. Haykin, Neural Networks and Learning Machine, (3rd Edition), Pearson 2008
 - Chapter 1




- Christopher M. Bishop, Pattern Recognition and Machine Learning (Information Science and Statistics), Springer 2006
 - Chapter 4



- Hands-On Machine Learning with Scikit-Learn and TensorFlow, Aurélien Géron, O'Reilly Media; 1 edition, 2017
 - Chapter 4

Literature



- Machine Learning - A Journey to Deep Learning, A. Wichert, Luis Sa-Couto, World Scientific, 2021
 - Chapter 5